

Name: Ngô Lê Ngọc Quý

Student ID: 21040003

Course: Artificial Intelligence

Lecturer: Trần Thế Vũ

CHALLENGE 3 REPORT

I. Overview.....	2
II. Phases.....	2
1. Data Collection and Preprocessing.....	2
2. Model Building and Training	2
3. Model Predicting.....	14
4. Deployment.....	17
4.1. Creating Template	17
4.2. Using Flask server	20
III. Conclusion.....	21
IV. Reference	22

I. Overview

The primary objective of this challenge is to create a robust system dedicated to the identification and prevention of violent, offensive, and harmful language on the internet. The project consists of four main phases: data collection and preprocessing, model building and training, prediction and deployment.

Github repository of the challenge: [<https://github.com/ngocquyxs/-AI-Challenge-3.git>]

II. Phases

1. *Data Collection and Preprocessing*

This code loads a dataset, explores its structure, visualizes the distribution of labels, and then splits it into training and testing sets for further machine learning model development.

```
# Load dataset

%cd /content/drive/MyDrive/[AI] Ngoc Quy/Bully

data = pd.read_csv('data/data.csv')

data = data.iloc[:, 1:]

print(data.info())

data.head()

# Visualize the distribution of labels

sns.countplot(x='label', data=data)

# Create training dataset

train_df = data.iloc[:int(len(data)* 0.8)]

train_df.info()

# Create testing dataset

test_df = data.iloc[len(train_df):]

test_df.info()
```

2. *Model Building and Training*

This code is setting up stratified k-fold cross-validation for the training dataset (train_df).

```
# "StratifiedKFold" class
```

```
skf = StratifiedKFold(n_splits=N_SPLITS)
```

```
for fold, (_, val_) in enumerate(skf.split(X=train_df, y=train_df.label)):
```

```
    train_df.loc[val_, "kfold"] = fold
```

This code is using the Hugging Face transformers library to load a tokenizer for the PhoBERT (Pre-trained models for Vietnamese) model.

```
# Load a tokenizer for the PhoBERT
```

```
tokenizer = AutoTokenizer.from_pretrained("vinai/phobert-base", use_fast=False)
```

This code defines a custom PyTorch dataset class called “SentimentDataset” for sentiment analysis.

```
class SentimentDataset(Dataset):
```

```
    # Initialization
```

```
    def __init__(self, df, tokenizer, max_len=120):
```

```
        self.df = df
```

```
        self.max_len = max_len
```

```
        self.tokenizer = tokenizer
```

```
    # Length Method
```

```
    def __len__(self):
```

```
        return len(self.df)
```

```
    # Get Item Method
```

```
    def __getitem__(self, index):
```

```
        row = self.df.iloc[index]
```

```
text, label = self.get_input_data(row)

encoding = self.tokenizer.encode_plus(
    text,
    truncation=True,
    add_special_tokens=True,
    max_length=self.max_len,
    padding='max_length',
    return_attention_mask=True,
    return_token_type_ids=False,
    return_tensors='pt',
)

return {
    'text': text,
    'input_ids': encoding['input_ids'].flatten(),
    'attention_masks': encoding['attention_mask'].flatten(),
    'targets': torch.tensor(label, dtype=torch.long),
}
```

Get Input Data Method

```
def get_input_data(self, row):
    # Preprocessing: {remove icon, special character, lower}
    text = row['content']
    text = ''.join(simple_preprocess(text))
    label = row['label']
```

```
return text, label
```

This code is generating a distribution plot of the token count in sentences to provide insights into the length distribution of the text data.

Distribution of length of Sentence

```
all_data = train_df.content.tolist() + test_df.content.tolist()

all_data = [' '.join(simple_preprocess(text)) for text in all_data]

encoded_text = [tokenizer.encode(text, add_special_tokens=True) for text in all_data]

token_lens = [len(text) for text in encoded_text]

sns.displot(token_lens)

plt.xlim([0,max(token_lens)])

plt.xlabel('Token Count')
```

This model is designed for sentiment classification tasks, where the PhoBERT model is used as a feature extractor, and a linear layer is employed for class prediction.

```
class SentimentClassifier(nn.Module):

    # Initialization

    def __init__(self, n_classes, device):

        super(SentimentClassifier, self).__init__()

        # Model Architecture

        self.bert = AutoModel.from_pretrained("vinai/phobert-base")

        self.drop = nn.Dropout(p=0.3)

        self.fc = nn.Linear(self.bert.config.hidden_size, n_classes)

        nn.init.normal_(self.fc.weight, std=0.02)

        nn.init.normal_(self.fc.bias, 0)
```

```
self.device = device
```

Forward Method

```
def forward(self, input_ids, attention_mask):
```

```
    last_hidden_state, output = self.bert(
```

```
        input_ids=input_ids,
```

```
        attention_mask=attention_mask,
```

```
        return_dict=False
```

```
    )
```

```
    x = self.drop(output)
```

```
    x = self.fc(x)
```

```
    return x
```

This code defines two functions: train for training the model and eval for evaluating the model on either the validation or test set.

"train" function

```
def train(model, criterion, optimizer, train_loader):
```

```
    model.train()
```

```
    losses = []
```

```
    correct = 0
```

```
    for data in train_loader:
```

```
        input_ids = data['input_ids'].to(device)
```

```
        attention_mask = data['attention_masks'].to(device)
```

```
        targets = data['targets'].to(device)
```

```
optimizer.zero_grad()

outputs = model(
    input_ids=input_ids,
    attention_mask=attention_mask
)

loss = criterion(outputs, targets)
_, pred = torch.max(outputs, dim=1)

correct += torch.sum(pred == targets)
losses.append(loss.item())
loss.backward()
nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
optimizer.step()
lr_scheduler.step()

print(f'Train Accuracy: {correct.double()/len(train_loader.dataset)} Loss:
{np.mean(losses)}')

# "eval" function
def eval(test_data = False):
    model.eval()
    losses = []
    correct = 0

    with torch.no_grad():
```

```
data_loader = test_loader if test_data else valid_loader

for data in data_loader:

    input_ids = data['input_ids'].to(device)

    attention_mask = data['attention_masks'].to(device)

    targets = data['targets'].to(device)


    outputs = model(

        input_ids=input_ids,

        attention_mask=attention_mask

    )


    _, pred = torch.max(outputs, dim=1)


    loss = criterion(outputs, targets)

    correct += torch.sum(pred == targets)

    losses.append(loss.item())


if test_data:

    print(f'Test Accuracy: {correct.double()/len(test_loader.dataset)} Loss:
{np.mean(losses)}')

    return correct.double()/len(test_loader.dataset)

else:

    print(f'Valid Accuracy: {correct.double()/len(valid_loader.dataset)} Loss:
{np.mean(losses)}')

    return correct.double()/len(valid_loader.dataset)
```


This code defines a function `prepare_loaders` to create PyTorch DataLoader objects for training and validation datasets.

```
# "prepare_loaders" function

def prepare_loaders(df, fold):

    df_train = df[df.kfold != fold].reset_index(drop=True)

    df_valid = df[df.kfold == fold].reset_index(drop=True)

    train_dataset = SentimentDataset(df_train, tokenizer, max_len=120)

    valid_dataset = SentimentDataset(df_valid, tokenizer, max_len=120)

    train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True, num_workers=2)

    valid_loader = DataLoader(valid_dataset, batch_size=16, shuffle=True, num_workers=2)

    return train_loader, valid_loader
```

This code is performing training and evaluation in a cross-validation loop for a sentiment analysis model.

```
# Training and evaluation

for fold in range(skf.n_splits):

    print(f'-----Fold: {fold+1} -----')

    train_loader, valid_loader = prepare_loaders(train_df, fold=fold)

    model = SentimentClassifier(n_classes=7, device = device).to(device)

    criterion = nn.CrossEntropyLoss()

    # Recommendation by BERT: lr: 5e-5, 2e-5, 3e-5

    # Batchsize: 16, 32

    optimizer = AdamW(model.parameters(), lr=2e-5)
```

```
lr_scheduler = get_linear_schedule_with_warmup(
    optimizer,
    num_warmup_steps=0,
    num_training_steps=len(train_loader)*EPOCHS
)

best_acc = 0

for epoch in range(EPOCHS):
    print(f'Epoch {epoch+1}/{EPOCHS}')
    print('-'*30)

    train(model, criterion, optimizer, train_loader)
    val_acc = eval()

    if val_acc > best_acc:
        torch.save(model.state_dict(), f'phobert_fold{fold+1}.pth')
        best_acc = val_acc
```

This code defines a test function for evaluating the ensemble model on a test dataset.

"test" function

```
def test(data_loader):
    models = []

    for fold in range(skf.n_splits):
        model = SentimentClassifier(n_classes=7, device = device)
        model.to(device)
        model.load_state_dict(torch.load(f'phobert_fold{fold+1}.pth'))
```

```
model.eval()

models.append(model)

texts = []

predicts = []

predict_probs = []

real_values = []

for data in data_loader:

    text = data['text']

    input_ids = data['input_ids'].to(device)

    attention_mask = data['attention_masks'].to(device)

    targets = data['targets'].to(device)

    total_outs = []

    for model in models:

        with torch.no_grad():

            outputs = model(

                input_ids=input_ids,

                attention_mask=attention_mask

            )

            total_outs.append(outputs)

    total_outs = torch.stack(total_outs)

    _, pred = torch.max(total_outs.mean(0), dim=1)

    texts.extend(text)
```

```

predicts.extend(pred)

predict_probs.extend(total_outs.mean(0))

real_values.extend(targets)


predicts = torch.stack(predicts).cpu()
predict_probs = torch.stack(predict_probs).cpu()
real_values = torch.stack(real_values).cpu()

print(classification_report(real_values, predicts))

return real_values, predicts

```

This code is testing sentiment analysis model on the test dataset.

```

test_dataset = SentimentDataset(test_df, tokenizer, max_len=50)

test_loader = DataLoader(test_dataset, batch_size=16, shuffle=True, num_workers=2)

real_values, predicts = test(test_loader)

```

	precision	recall	f1-score	support
0	0.89	0.89	0.89	1144
1	0.88	0.89	0.88	1068
accuracy			0.89	2212
macro avg	0.89	0.89	0.89	2212
weighted avg	0.89	0.89	0.89	2212

This code is using Seaborn to create a heatmap of the confusion matrix for sentiment analysis model's predictions on the test dataset.

```

class_names = [0, 1]

sns.heatmap(confusion_matrix(real_values, predicts), annot=False, xticklabels =
class_names, yticklabels = class_names)

```

This code is implemented a function “check_wrong” to identify and print some examples where sentiment analysis model made incorrect predictions.

```
# "check_wrong" function
```

```
def check_wrong(real_values, predicts):
```

```
    wrong_arr = []
```

```
    wrong_label = []
```

```
    for i in range(len(predicts)):
```

```
        if predicts[i] != real_values[i]:
```

```
            wrong_arr.append(i)
```

```
            wrong_label.append(predicts[i])
```

```
    return wrong_arr, wrong_label
```

```
for i in range(15):
```

```
    print('-'*50)
```

```
    wrong_arr, wrong_label = check_wrong(real_values, predicts)
```

```
    print(test_df.iloc[wrong_arr[i]].content)
```

```
    print(f'Predicted: ({class_names[wrong_label[i]]) --vs-- Real label:  
({class_names[real_values[wrong_arr[i]]])'})
```

This code is implementing an “infer” function for making predictions on a single input text using sentiment analysis model.

```
# "infer" function
```

```
def infer(text, tokenizer, model, max_len=120):
```

```
    encoded_review = tokenizer.encode_plus(  
        text,  
        max_length=max_len,  
        truncation=True,
```

```
add_special_tokens=True,  
padding='max_length',  
return_attention_mask=True,  
return_token_type_ids=False,  
return_tensors='pt',  
)  
  
input_ids = encoded_review['input_ids'].to(device)  
attention_mask = encoded_review['attention_mask'].to(device)  
  
output = model(input_ids, attention_mask)  
_, y_pred = torch.max(output, dim=1)  
  
print(f'Text: {text}')print(f'Negative: {class_names[y_pred]}')
```

This code is saving model.

Save model

```
output_path = "/content/drive/MyDrive/[AI] Ngoc Quy/Bully/bully_detection_model"  
torch.save(model.state_dict(), f'{output_path}.pth')  
tokenizer.save_pretrained(output_path)
```

3. Model Predicting

This code is

"load_model_and_tokenizer" function

```
def load_model_and_tokenizer(model_class, tokenizer_class, model_path, tokenizer_path,
```

device):

Load model

```
model = model_class(n_classes=7, device=device)

model.load_state_dict(torch.load(model_path, map_location=torch.device('cpu')))

model.eval()
```

Load tokenizer

```
tokenizer = tokenizer_class.from_pretrained(tokenizer_path)
```

```
return model, tokenizer
```

```
model_path = '/content/drive/MyDrive/[AI] Ngoc Quy/Bully/bully_detection_model.pth'
```

```
tokenizer_path = '/content/drive/MyDrive/[AI] Ngoc Quy/Bully/bully_detection_model'
```

Load saved model and tokenizer

```
loaded_model, loaded_tokenizer = load_model_and_tokenizer(SentimentClassifier,
AutoTokenizer, model_path, tokenizer_path, device)
```

This code is making sentiment predictions on a given input sentence using a loaded sentiment analysis model and tokenizer.

```
class_names = [0, 1]
```

"predict_sentiment" function

```
def predict_sentiment(sentence, tokenizer, model, device, max_len=120):
```

Tokenize the input sentence

```
    encoded_input = tokenizer.encode_plus(
        sentence,
        max_length=max_len,
```

```
truncation=True,  
add_special_tokens=True,  
padding='max_length',  
return_attention_mask=True,  
return_token_type_ids=False,  
return_tensors='pt',  
)  
  
# Move input tensors to the appropriate device  
input_ids = encoded_input['input_ids'].to(device)  
attention_mask = encoded_input['attention_mask'].to(device)
```

Make the prediction

```
with torch.no_grad():  
    output = model(input_ids, attention_mask)  
    _, predicted_label = torch.max(output, dim=1)  
  
return predicted_label.item()
```

Example usage

```
sentence_to_predict = "Chào buổi sáng"  
  
predicted_label = predict_sentiment(sentence_to_predict, loaded_tokenizer, loaded_model,  
device)  
  
print(f"Predicted Sentiment Label: {class_names[predicted_label]}")
```


4. Deployment

4.1. Creating Template

This HTML template is a crucial part of web application, as it allows users to interact with sentiment analysis model by entering sentences and viewing the predictions on the web page.

```
<!DOCTYPE html>

<html lang="en">

<head>

  <meta charset="UTF-8">

  <meta name="viewport" content="width=device-width, initial-scale=1.0">

  <title>Sentiment Analysis</title>

  <style>

    body {

      font-family: 'Arial', sans-serif;

      margin: 20px;

      text-align: center;

      background-color: #f4f4f4;

    }

    h1 {

      color: #333;

    }

    form {

      max-width: 400px;

      margin: auto;

      background-color: #fff;

      padding: 20px;

      border-radius: 8px;
```

```
    box-shadow: 0 0 10px rgba(0, 0, 0, 0.1);
  }

  label {
    display: block;
    margin-top: 10px;
    color: #555;
    font-weight: bold;
  }

  input[type="text"] {
    width: calc(100% - 22px);
    padding: 10px;
    margin-top: 5px;
    margin-bottom: 20px;
    box-sizing: border-box;
    border: 1px solid #ccc;
    border-radius: 4px;
    display: inline-block;
  }

  button {
    background-color: #4CAF50;
    color: white;
    padding: 10px 15px;
    border: none;
    border-radius: 4px;
    cursor: pointer;
    font-size: 16px;
```

```
}

button:hover {

    background-color: #45a049;

}

h2 {

    margin-top: 20px;

    color: #333;

}

p#result {

    font-size: 18px;

    font-weight: bold;

    color: #333;

}

</style>
</head>
<body>

<h1>Sentiment Analysis</h1>

<form action="/predict" method="post">

    <label for="sentence">Enter a sentence:</label>

    <input type="text" name="sentence" id="sentence" required>

    <br>

    <button type="submit">Predict</button>

</form>

<h2>Result:</h2>

<p id="result">{{ result }}</p>

</body>
```

</html>

4.2. Using Flask server

This Flask application loads a sentiment analysis model and tokenizer, provides a web interface for users to input sentences, predicts the sentiment of the input sentences, and displays the results on the web page.

```
from flask import Flask, render_template, request

import os

import torch

import speech_recognition as sr

from transformers import AutoTokenizer

from model import SentimentClassifier, load_model_and_tokenizer, predict_sentiment

app = Flask(__name__)

# Load saved model and tokenizer

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

model_path = r'E:\ngocquy_python\AI\bully\bully_detection_model.pth'

tokenizer_path = r'E:\ngocquy_python\AI\bully\bully_detection_model'

loaded_model, loaded_tokenizer = load_model_and_tokenizer(SentimentClassifier,
AutoTokenizer, model_path, tokenizer_path, device)

class_names = {0: 'Non-Violence', 1: 'Violence'} # Update with your actual class names

@app.route('/')

def index():

    return render_template('index.html')
```

```
@app.route('/predict', methods=['POST'])
def predict():
    if request.method == 'POST':
        sentence = request.form['sentence']
        predicted_label = predict_sentiment(sentence, loaded_tokenizer, loaded_model, device)

        result = class_names[predicted_label]
        return render_template('index.html', result=result)

if __name__ == '__main__':
    app.run(debug=True)
```

III. Conclusion

This challenge has successfully implemented development of a comprehensive system for sentiment analysis and the deployment of Flask-based web application to make predictions based on a pre-trained model.

Sentiment Analysis

Enter a sentence:

Result:

Non-Violence

IV. Reference

The above challenge has reference from

[https://github.com/phusroyal/ViHOS?fbclid=IwAR2rAnYcXjgj4YQ4WXthjYJrb_ij_hwLbz2sF3_Y5lJierEa7VOyT3Aw5fw]

[<https://www.kaggle.com/code/trnmtin/phobert-classification-for-vietnamese-text/notebook?fbclid=IwAR32To1PbBtC7eWpS2cglKWhzg78balFHpK5CcYwyBztXzfUBnCwjpFNIk>]